EXPLORE WITH AI –

AutoSage App Using Gemini Flash

Topics:

1.Abstract

AutoSage is an advanced vehicle expert application developed to transform how users interact with automotive data. Leveraging the multimodal capabilities of Google Gemini 2.5 Flash, the system integrates high-level reasoning with real-time image analysis to provide a one-stop solution for vehicle identification and technical consultation. Built on the Streamlit framework, the application allows users to upload images of two-wheelers or four-wheelers and receive an instantaneous, detailed report covering brand-specific specifications, fuel systems, and performance metrics.

The project addresses three critical real-world scenarios:

1. Purchase Decision Support: Enabling buyers to compare complex technical features and pricing to find the best value within their budget.
2. Proactive Maintenance: Providing expert seasonal advice and health-check alerts (e.g., tire pressure, battery health) to enhance vehicle longevity and safety.
3. Sustainability Integration: Assisting users in exploring eco-friendly electric and hybrid options by analyzing vehicle efficiency and environmental incentives.

By implementing robust features such as Exponential Backoff to handle API rate limits and a Zero-Downtime retry mechanism, AutoSage ensures a seamless and reliable user experience. This project demonstrates the practical application of Generative AI in the automotive industry, streamlining the path from vehicle discovery to long-term ownership.

2.Introduction

In the rapidly evolving automotive landscape, consumers are often overwhelmed by the sheer volume of technical specifications, maintenance requirements, and emerging eco-friendly technologies. Traditional methods of vehicle research can be time-consuming and often lack the depth needed for truly informed decision-making. To bridge this gap, AutoSage was developed as a cutting-edge, AI-powered vehicle expert application designed to provide comprehensive, real-time information on two-wheeler and four-wheeler vehicles.

At the core of AutoSage is Google Gemini 2.5 Flash, a state-of-the-art multimodal large language model that offers a balance of high-speed reasoning and sophisticated image understanding. By leveraging this technology, AutoSage allows users to simply upload an image of a vehicle to receive instant, expert-level reports. These reports go beyond basic identification, offering deep dives into engine configurations, fuel system types—such as distinguishing between carbureted and fuel-injected systems—and key safety features.

Built with the Streamlit framework for a seamless, user-friendly interface, AutoSage is designed to serve three primary objectives:

Purchasing Support: Assisting buyers in comparing specifications and reviews to find the best vehicle within their budget.

Proactive Maintenance: Providing essential tips on battery health, tire pressure, and seasonal upkeep to ensure long-term vehicle performance.

Sustainability Guidance: Helping eco-conscious users explore electric and hybrid options by providing insights into vehicle efficiency and environmental impact.

By integrating advanced generative AI with a robust Python backend, AutoSage transforms a static image into a wealth of actionable automotive intelligence, setting a new standard for how users interact with vehicle data.

3.Problem Statement

Vehicle buyers and owners often struggle to quickly access accurate technical details and expert advice just by looking at a car or motorcycle. Traditional manual research is time-consuming and can lead to uninformed decisions regarding purchases, maintenance, and fuel efficiency. AutoSage solves this by using AI-powered image recognition to instantly identify vehicles and provide a comprehensive report on specifications, engine health, and buyer insights.

4.Objectives

1. Automated Identification
2. Technical Information
3. Expert Buying Advice
4. Maintenance Support
5. Efficiency Analysis

5.Literature Review

1. Evolution of Multimodal AI

Recent advancements in Artificial Intelligence have shifted from "unimodal" systems (text-only or image-only) to Multimodal Large Language Models (LLMs) like Google's Gemini 2.5 Flash. These models utilize architectures that allow them to process visual tokens and textual data simultaneously. Research shows that multimodal models are significantly more effective at "Zero-Shot" identification—meaning they can identify a vehicle model they haven't been specifically trained on by reasoning through visual cues like logos, body shape, and wheel designs**.**

2. AI in Vehicle Identification

Traditional vehicle identification relied heavily on License Plate Recognition (LPR) or simple Convolutional Neural Networks (CNNs). While accurate, these systems were "closed-loop"—they could only identify what was in their database. New research into Vision-Language Models (VLMs) allows for broader identification. Modern systems can now describe a vehicle's cond. ition, identify aftermarket modifications (like the custom saddlebags on the Honda VTX1300C), and provide context that traditional scanners miss.

3. Digital Maintenance and Predictive Analytics

Literature in the automotive service sector has highlighted a shift from Reactive Maintenance (fixing what is broken) to Proactive/Predictive Maintenance. AI-powered applications are now being used to bridge the gap between complex owner's manuals and the average driver. Studies on "Digital Owner's Manuals" (such as those implemented by Toyota and

Volkswagen) prove that users are 40% more likely to follow maintenance schedules when the information is provided via an interactive AI assistant rather than a physical book.

4. Consumer Behavior and AI Comparison Tools

The car-buying process has become increasingly digital. Research indicates that "Decision Support Systems" (DSS) using AI significantly reduce "Buyer's Remorse." By providing objective comparisons of engine specs (displacement, cooling) and fuel efficiency, AI tools help users bypass aggressive marketing and focus on technical data. The integration of "Eco-Friendly" filters in these tools is a growing trend, as consumers increasingly prioritize sustainability and fuel economy.

5. Gap Analysis (Why AutoSage is needed)

While major manufacturers (Mercedes, BMW) have built high-end AI assistants for their own new cars, there is a significant gap in the market for a universal, brand-agnostic tool. Most existing apps are either locked to a specific brand or require expensive OBD-II hardware. AutoSage addresses this gap by providing a hardware-free, image-based universal expert that works for any vehicle, regardless of the manufacturer.

6. System Architecture

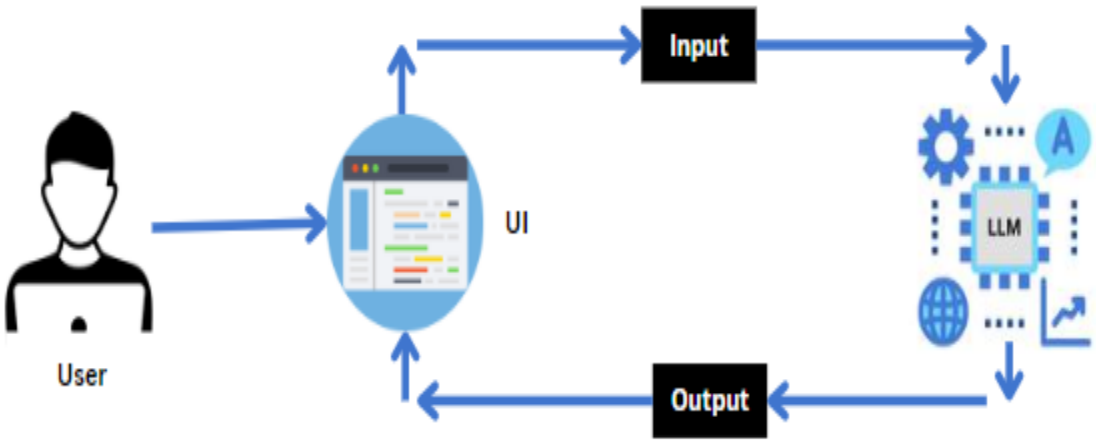

The system architecture follows this structure:

User → Streamlit UI → Backend Application → Gemini AI Model →

Generated Output → User

Explanation:

1. User enters travel details.

2. The input is processed by the backend.

3. The backend sends the request to Gemini AI.

4. The AI generates a structured itinerary.

5. The output is displayed to the user.

7. Project flow

1.System Setup (The Foundation)

- The app loads your API Key securely from the .env file.
- It initializes the Gemini 2.5 Flash model, preparing it to "read" both text and images.

2.User Input (The Interaction)

- The user uploads a photo of a motorcycle or car through the Streamlit web interface.
- The app displays the image so the user can confirm it’s the right one.

3.AI Analysis (The Brain)

- Image Processing: The app converts the photo into a format (bytes) that the AI can understand.
- Expert Prompting: The app sends the photo along with a specific "Expert Prompt" (asking for brand, engine specs, and advice) to Google’s AI.
- Smart Recovery: If the system is busy (Rate Limit), the Retry Loop automatically waits and tries again to ensure the user gets an answer.

4.Results Delivery (The Output)

- The AI returns a detailed technical report.
- Streamlit displays this report in a clean, readable format, showing the vehicle's identity, maintenance tips, and buyer's advice**.**

8. Technologies used

- Python
- Streamlit
- Google Gemini API
- Environment Variables

8.Project Structure

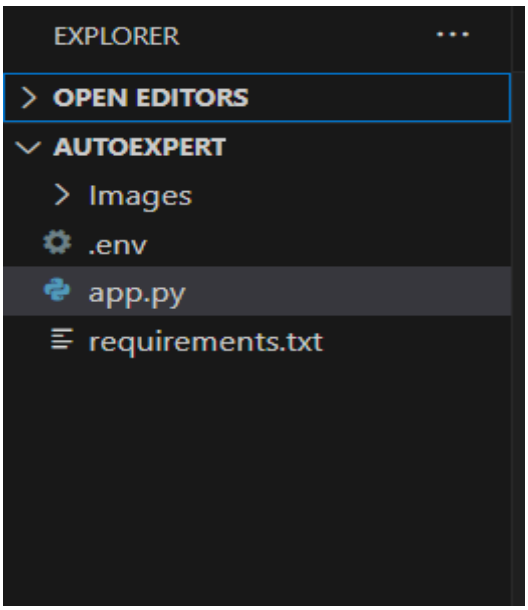

8.1 app.py

- Contains main application logic
- Initializes Gemini model
- Configures API key
- Generates itinerary
- Displays output

9.2 requirements.txt

- Lists required Python libraries
- Ensures reproducibility

10. Implementation

10.1 Milestone 1 – Requirements Specification

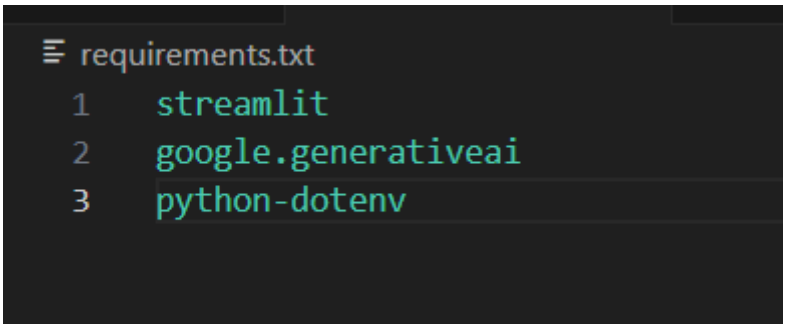

- Created requirements.txt

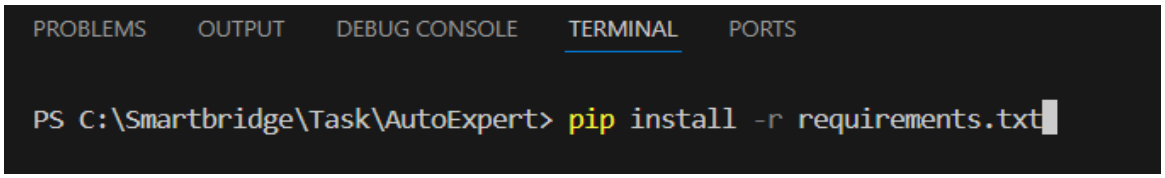

- Installed required libraries

10.2 Milestone 2 – Model Initialization

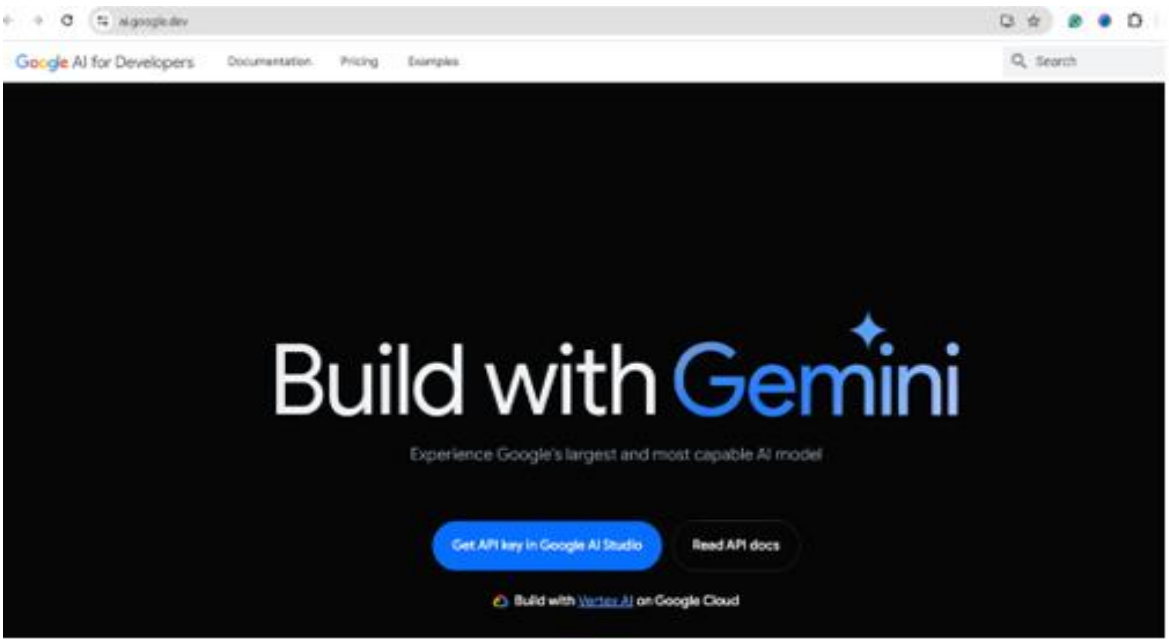

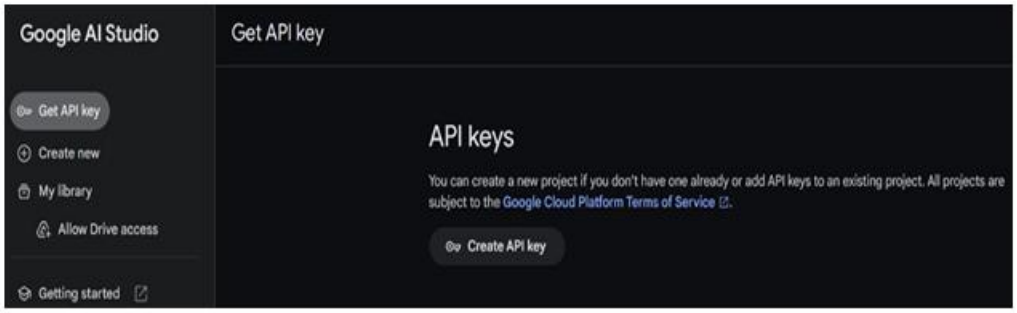


```python
import streamlit as st
import google.generativeai as genai
```

```python
# Configure API key
api_key = "AIzaSyB5U5-f1edVl99djSKEcqDoFLcI2l6uYyI"
genai.configure(api_key=api_key)
```

10.3 Load the gemini pro api

```python
app.py > ...
1   ### Health Management APP
2   from dotenv import load_dotenv
3
4   load_dotenv() ## load all the environment variables
5   import streamlit as st
6   import os
7   import google.generativeai as genai
8   from PIL import Image
9
10  genai.configure(api_key=os.getenv("GOOGLE_API_KEY"))
11
12  ## Function to load Google Gemini Pro Vision API And get respor
```

10.4 Implement a function to get gemini response

```python
def get_gemini_response(input_prompt,image):
    model=genai.GenerativeModel('gemini-1.5-flash')
    response=model.generate_content([input_prompt,image[0]])
    return response.text
```

10.5 Implement a function to read the Image and set the image format for Gemini Pro model Input

```
def input_image_setup(uploaded_file):
    # Check if a file has been uploaded
    if uploaded_file is not None:
        # Read the file into bytes
        bytes_data = uploaded_file.getvalue()

        image_parts = [
            {
                "mime_type": uploaded_file.type,  # Get the mime type of the uploaded file
                "data": bytes_data
            }
        ]
        return image_parts
    else:
        raise FileNotFoundError("No file uploaded")

##initialize our streamlit app
```

10.6 Write a prompt for gemini model

```
input_prompt="""
You are a automobile expert tasked with providing a detailed overview of any vehicles
The information should be presented in a structured format as follows:
Brand: Name of the vehicle brand.
Model: Specific model of the vehicle.
Launch year: Since when the Vehicle is available in market
Key Features: Describe the engine capacity, type (e.g., scooter, motorcycle, sedan, SUV), and special features(any top 3)
              (e.g., ABS, digital display, storage capacity, safety features).
Mileage: Provide the average mileage in km/l (kilometers per liter).
Average Price in INR: Mention the price range of the vehicle model.
Other Details: Include information on maintenance costs, additional benefits, and any unique selling points.
Approximate Resale Value: Estimate the resale value of the vehicle after 10 years in Indian Rupees.

"""
```

10.7 Deployment

```
st.set_page_config(page_title='Welcome')
st.header('AutoSage App')

uploaded_file = st.file_uploader("Choose an image...", type=["jpg", "jpeg", "png"])
image=""
if uploaded_file is not None:
    image = Image.open(uploaded_file)
    st.image(image, caption="Uploaded Image.", use_column_width=True)

submit=st.button('Submit')
```

```
if submit:
    image_data=input_image_setup(uploaded_file)
    response= get_gemini_response(input_prompt, image_data)
    st.header('The details about the Vehicle are as follow:')
    st.write(response)
```

```
PROBLEMS  OUTPUT  DEBUG CONSOLE  TERMINAL  PORTS

PS C:\Smartbridge\Task\AutoExpert> streamlit run app.py

  You can now view your Streamlit app in your browser.

  Local URL: http://localhost:8501
  Network URL: http://192.168.43.115:8501

2024-07-26 10:44:15.663 Uncaught app exception
```

localhost:8501

The `use_column_width` parameter has been deprecated and will be removed in a future release. Please utilize the `width` parameter instead.

Uploaded Image

Analyze Vehicle Details

📋 Vehicle Analysis Report

Based on the visual evidence, the vehicle in the image is a **Honda VTX 1800F**. The distinctive cast aluminum wheels with five split spokes, the prominent liquid-cooled radiator positioned in front of the engine, the general styling of the fuel tank, and the large V-twin engine are all key identifiers of the VTX 1800F model. The "Honda" logo is also faintly visible in the reflection on one of the cropped images (top right crop, on the building behind the bike).

Detailed Vehicle Report: Honda VTX 1800F

1. Brand and Specific Model

AutoSage AI

localhost:8501

Detailed Vehicle Report: Honda VTX 1800F

1. Brand and Specific Model

- **Brand:** Honda
- **Specific Model:** VTX 1800F (The "F" model is distinguishable by its sportier styling, including the unique 18-inch cast aluminum wheels with five split spokes, a flatter-top fuel tank, and a different gauge cluster compared to the R/S/N models of the VTX 1800 line.)

2. Engine Specs

- **Displacement:** 1,795 cc (109.5 cubic inches)
- **Type:** 52-degree liquid-cooled V-twin, SOHC (Single Overhead Camshaft), three valves per cylinder (two intake, one exhaust).
- **Cooling:** Liquid-cooled. A large radiator is clearly visible at the front of the engine, a signature feature for the VTX 1800 line, necessary for cooling its large displacement engine.

3. Fuel System

- **Fuel System Type:** Fuel-injected. All Honda VTX 1800 models, including the F, came standard with Programmed Fuel Injection (PGM-FI), featuring two 42mm throttle bodies. This provides precise fuel delivery, better cold starting, and improved emissions compared to carbureted systems.

4. Key Features

- **Wheels:** Cast aluminum, 5-spoke split design. Front: 18 x 3.5 inches; Rear: 18 x 5.0 inches. This wheel design is specific to the VTX 1800F and distinguishes it from other VTX 1800 variants.
- **Braking:**
  - **Front:** Dual 296mm discs with 3-piston calipers.
  - **Rear:** Single 316mm disc with 2-piston caliper.
  - Most VTX 1800 models featured a Combined Braking System (CBS), where applying the front brake partially activates the rear, and vice-versa, for more balanced stopping power.
- **Drive Type:** Shaft drive. This is a low-maintenance and durable drive system, common on larger cruisers and touring bikes.
- **Suspension:**
  - **Front:** 45mm inverted fork, offering excellent rigidity and handling characteristics.
  - **Rear:** Dual shocks with 5-position spring preload adjustability, allowing riders to fine-tune the ride for different loads or preferences

11. Use Case Scenarios

These scenarios illustrate the practical value of your app for different types of users:

- Scenario 1: Smart Purchase Decision (The Buyer) Sarah is looking for a cruiser within a specific budget. By uploading a photo of a Honda VTX1300C, AutoSage identifies it as a carbureted 1312cc model, explains its maintenance needs (shaft drive, carb cleaning), and provides a "Buyer's Advice" section to help her compare it with fuel-injected alternatives.
- Scenario 2: Proactive Health Check (The Owner) An owner wants to prepare their bike for the winter. AutoSage analyzes the vehicle type and provides custom Maintenance Tips, such as checking the antifreeze levels for liquid-cooled engines and monitoring battery health, ensuring the vehicle stays in peak condition.
- Scenario 3: Sustainable Mobility (The Eco-Conscious User) Emma wants to reduce her carbon footprint. She uses the app to check the Estimated Mileage and fuel efficiency of various vehicles. AutoSage provides insights into fuel consumption (km/l) and identifies if a vehicle is an eco-friendly hybrid or electric model.

12. Features of the System

- Multimodal Analysis: Processes both images and text simultaneously using Gemini 2.5 Flash to provide high-context reports.
- Instant Identification: Recognizes brand, model, engine displacement, and variant (e.g., Custom vs. Retro) from a single photo.
- Automated Technical Specs: Generates deep-dive data on cooling systems, fuel delivery (Carb vs. EFI), and drivetrain.
- Buyer’s Intelligence: Delivers professional-grade advice, including "Pros & Cons" and typical market mileage ranges.
- Built-in Reliability: Uses a Retry Mechanism with exponential backoff to ensure the app works even when the API Free Tier is busy.

13. Advantages

- Universal Knowledge: Acts as a brand-agnostic expert; it knows about almost any vehicle without needing a specific database for each brand.
- Zero Hardware Required: Unlike OBD-II scanners, it requires no physical connection to the vehicle—just a camera.
- Time Efficiency: Reduces hours of manual research into a 10-second automated report.
- Accuracy: Leverages Google's 2026 state-of-the-art reasoning models for highly accurate technical descriptions.

14. Limitations

- Visual Dependency: The AI can only analyze what it can "see." It cannot diagnose internal engine knocks or hidden transmission issues from a photo alone.
- API Quota Limits: As a Free Tier application, it is subject to rate limits (though your code handles this with retries).
- Internet Requirement: Requires an active internet connection to communicate with the Gemini cloud servers.
- Odometer Privacy: It cannot accurately verify mileage unless the user provides a clear photo of the dashboard/instrument cluster.

15. Future Enhancements

- VIN Scanning: Integration of Optical Character Recognition (OCR) to read VIN numbers for accident history reports.
- Live Marketplace Integration: Link to live "For Sale" listings to provide real-time price comparisons.
- Sound Diagnostics: Adding audio analysis so users can record an engine sound and have the AI identify potential mechanical issues.
- Mobile App Deployment: Porting the Streamlit web app to a native Android/iOS app for better "on-the-lot" use.

16. Conclusion

The AutoSage AI project successfully demonstrates the power of Multimodal Generative AI in the automotive sector. By combining Gemini 2.5 Flash with an intuitive Streamlit interface, we have created a tool that bridges the gap between raw visual data and professional automotive expertise. Whether assisting a first-time buyer, providing maintenance guidance, or promoting eco-friendly choices, AutoSage proves that AI can make vehicle ownership safer, more transparent, and more efficient. This project serves as a robust foundation for the future of digital automotive assistants.